

Next-Generation Automatic Number Plate Recognition (ANPR) System

1. What is this project?

The **Automatic Number Plate Recognition (ANPR) System** is a smart software application designed to automatically read license plates from images of vehicles. Think of it as a gatekeeper that has digital "eyes" and a "brain."

When you show it a picture of a car, it figures out where the license plate is, reads the letters and numbers on the plate, and checks a database to see if that specific vehicle is authorized (for example, to open a boom barrier at a parking lot or a toll booth).

2. How Does It Work?

The system acts like an assembly line, passing the image through several stages until it gets the final answer. Here is what happens behind the scenes in plain English:

Step 1: Image Upload (The "Eyes")

A user uploads an image of a car through a sleek, modern web interface. This interface securely sends the image to the "Backend" (our server) to start the analysis.

Step 2: Finding the Plate (Object Detection)

Before we can read the text, we have to find the plate itself. We use an Artificial Intelligence model called **YOLOv8** (You Only Look Once). This AI was specifically trained on thousands of images to identify bounding boxes around Indian license plates. It behaves like a spotlight, highlighting only the rectangular plate and ignoring the rest of the car, road, and background.

Step 3: Isolating the Letters (Segmentation)

Now that we have a cropped picture of just the license plate, we need to separate each individual character. We use computer vision techniques (like thresholding and contour detection) to slice the plate into multiple tiny images, where each tiny image contains exactly one letter or number.

Step 4: Reading the Text (Character Recognition)

Next, we pass those tiny sliced images to a second Artificial Intelligence model called a **Convolutional Neural Network (CNN)**. We trained this PyTorch AI to recognize the alphabet (A-Z) and numbers (0-9). It acts as the "Brain," predicting the character for each tiny image and combining them to form the final text string (e.g., "MH12AB1234").

Step 5: Authorization (The Database Check)

Finally, the system takes the predicted text string and checks it against a digital database of authorized vehicles. If the license plate matches a record in the database, the system grants access. If not, access is denied. Everything happens in a fraction of a second!

3. What Technology Powers This?

To build a system this robust, we combined several modern technologies:

- **The Frontend (User Interface):** Built with **React** and **Vite**, featuring a beautiful "glassmorphism" design. It gives the user a smooth, app-like experience to upload images and see the results instantly.
 - **The Backend (The Server):** Built with **FastAPI** (Python). It acts as the bridge that connects the user's web browser to our heavy Artificial Intelligence logic.
 - **The AI Engine:** Powered by **PyTorch** and **Ultralytics (YOLO)**. This is where the actual math and predictions happen, heavily utilizing the computer's Graphics Processing Unit (GPU) for lightning-fast speeds.
 - **The Database:** A lightweight **SQLite** database stores the list of vehicles allowed to pass.
-

4. Why is this useful?

This project can be deployed in the real world for:

- **Smart Parking Lots:** Automatically raising the barrier for monthly subscribers without needing a human guard or paper tickets.
 - **Toll Booths:** Automatically billing vehicles as they drive on the highway.
 - **Security & Law Enforcement:** Alerting authorities if a stolen or flagged vehicle drives past a camera.
-

5. Future Potential

While the system works beautifully right now, we can always make it smarter! Future updates could include taking a live video feed from a webcam to scan cars as they drive by in real-time, or adding an administrative dashboard to manage the database of allowed vehicles directly from the website.